

# Concurrency

---

Optimistic vs pessimistic approaches

**AMT 2018**

**Olivier Liechti**

# Introduction

# System overview: REST APIs



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
www.heig-vd.ch

The screenshot shows a GitHub repository page for 'SoftEng-HEIGVD / Teaching-HEIGVD-AMT-ConcurrentTransactions'. The repository is described as 'A demo project to experiment with concurrency issues and transactions in a Java EE application'. It has 1 commit, 8 branches, 0 releases, and 1 contributor. The latest commit is 9c2bf12 on Nov 9, 2015. The repository is licensed under MIT. A modal window for switching branches is open, showing a list of branches: master (selected), step1-validating-on-the-fly-account-creation, step2-really-validating-on-the-fly-account-creation, step3-fix-account-creation-with-try-catch, step4-fix-account-creation-with-upsert, step5-fix-account-creation-with-try-catch-pessimistic-lock, step6-fix-account-creation-with-try-catch-optimistic-lock, and step7-fix-account-creation-with-try-catch-optimistic-lock-retries. The repository name 'Teaching-HEIGVD-AMT-ConcurrentTransactions' is visible in the header. The description 'A demo project to experiment with concurrency issues and transactions in a Java EE application' is also visible. The repository is created by 'SoftEng-HEIGVD'.

SoftEng-HEIGVD / Teaching-HEIGVD-AMT-ConcurrentTransactions

Unwatch 2 Star 0 Fork 0

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

A demo project to experiment with concurrency issues and transactions in a Java EE application

Manage topics

1 commit 8 branches 0 releases 1 contributor MIT

Branch: master New pull request

Create new file Upload files Find file Clone or download

Switch branches/tags

Find or create a branch...

Branches Tags

- ✓ master
- step1-validating-on-the-fly-account-creation
- step2-really-validating-on-the-fly-account-creation
- step3-fix-account-creation-with-try-catch
- step4-fix-account-creation-with-upsert
- step5-fix-account-creation-with-try-catch-pessimistic-lock
- step6-fix-account-creation-with-try-catch-optimistic-lock
- step7-fix-account-creation-with-try-catch-optimistic-lock-retries

Latest commit 9c2bf12 on Nov 9, 2015

Initial commit	3 years ago
Initial commit	3 years ago
Initial commit	3 years ago

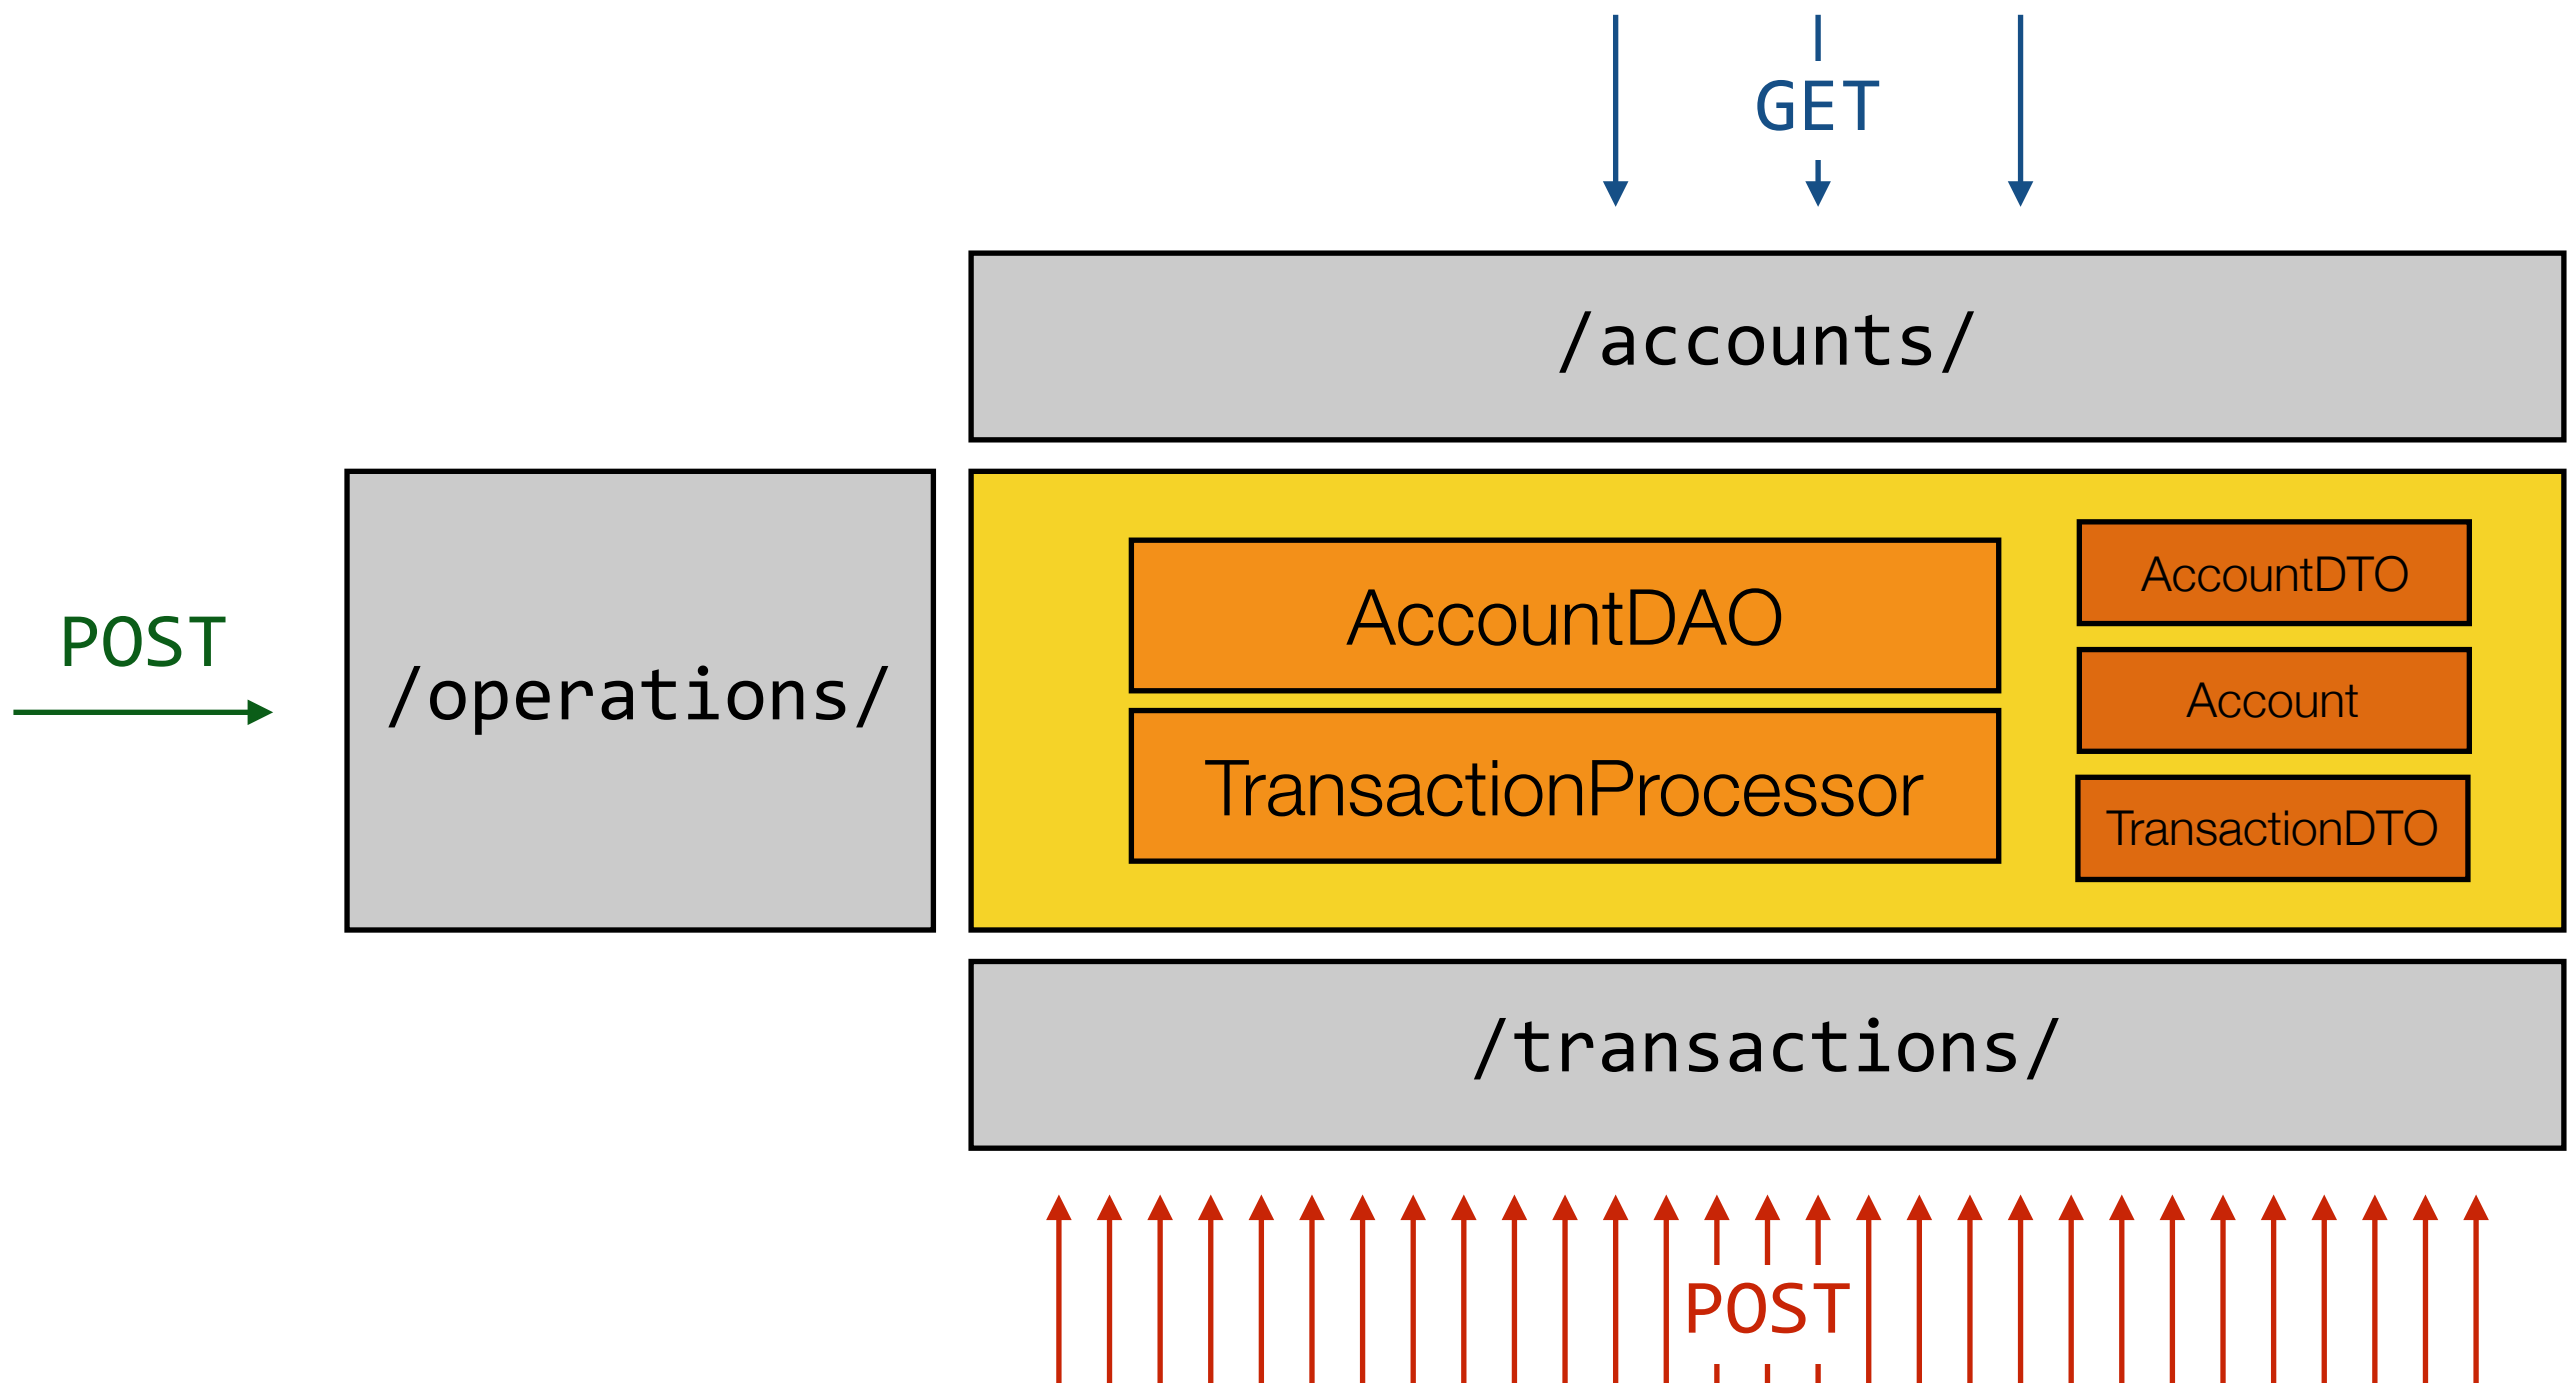
## Teaching-HEIGVD-AMT-ConcurrentTransactions

A demo project to experiment with concurrency issues and transactions in a Java EE application

Status Help

Contact GitHub Pricing API Training Blog About

# System overview: REST APIs



# System overview: REST APIs

---

## Account

```
id : long  
balance : double  
numberOfTransactions : long  
holderName : String
```

## Transaction

```
accountId : long  
amount : double
```

# Optimistic vs Pessimistic


# Transactions & concurrency control

- If several transactions are processed **concurrently**, unexpected results may occur. There are different strategies and mechanisms for dealing with that.

```
@Stateless
public class TransactionProcessor {

    @EJB
    AccountDAO accountDao;

    public void processTransaction(Transaction t) {
        Account a = accountDao.findById(t.getAccountId());
        long previousBalance = a.getBalance();
        a.setBalance(previousBalance + t.getAmount());
    }
}
```



What happens if another transaction modifies the account balance between these two statements?

# Optimistic concurrency control

---

- In many applications, there is a **high ratio of “read to write” operations** (many transactions read data, few update data). Moreover, there is a “small” likelihood that two concurrent transactions try to update the same data.
- In this case, for performance and scalability reasons, it is often recommended to implement an optimistic concurrency control mechanism.
- The mechanism works as follows:
  - When a program **reads** a record, it gets its “**version number**” (the number of previous updates) in a table column.
  - When it **updates** this record, it makes sure that the version number has not been incremented (this would indicate a conflict with another transaction).
- The developer has to write the logic to execute when a conflict is notified (retry, notify the user, etc.)



# Optimistic concurrency control with JPA

- **JPA supports optimistic concurrency control.**
- To use it, the first step is to annotate one field of the entity with the **@version** annotation. JPA will ensure that this value is incremented with every update.
- The second step is to catch the **OptimisticLockException** that may be thrown by JPA when the transaction commits.
- This is where the developer specifies what to do if a conflict has been detected. In some cases, it is possible to immediately and silently retry the transaction.

```
@Entity
public class Account {

    @Id
    long accountId;

    @Version
    long version;
}
```

# Pessimistic concurrency control

---

- When an optimistic concurrency control is not appropriate, then it is possible to implement **pessimistic concurrency control with locks**.
- RDBMS support different types of locks (read lock, write lock).
- When a transaction obtains a **read lock** on a record, it cannot be modified by other transactions. However, it can be read by other transactions.
- When a transaction obtains a **write lock** on a record, it cannot be modified, nor read by other transactions.
- Locking database records **introduce issues**: scalability, performance, deadlocks. It can be tricky to decide when to obtain a lock and for how long.

# Pessimistic concurrency control with JPA

- **JPA supports pessimistic concurrency control since version 2.0**
- It is possible to **lock a record** with **em.lock(entity, LOCK\_TYPE)**.
- It is also possible to lock a record at the time of retrieval with **em.find(class, id, LOCK\_TYPE)**

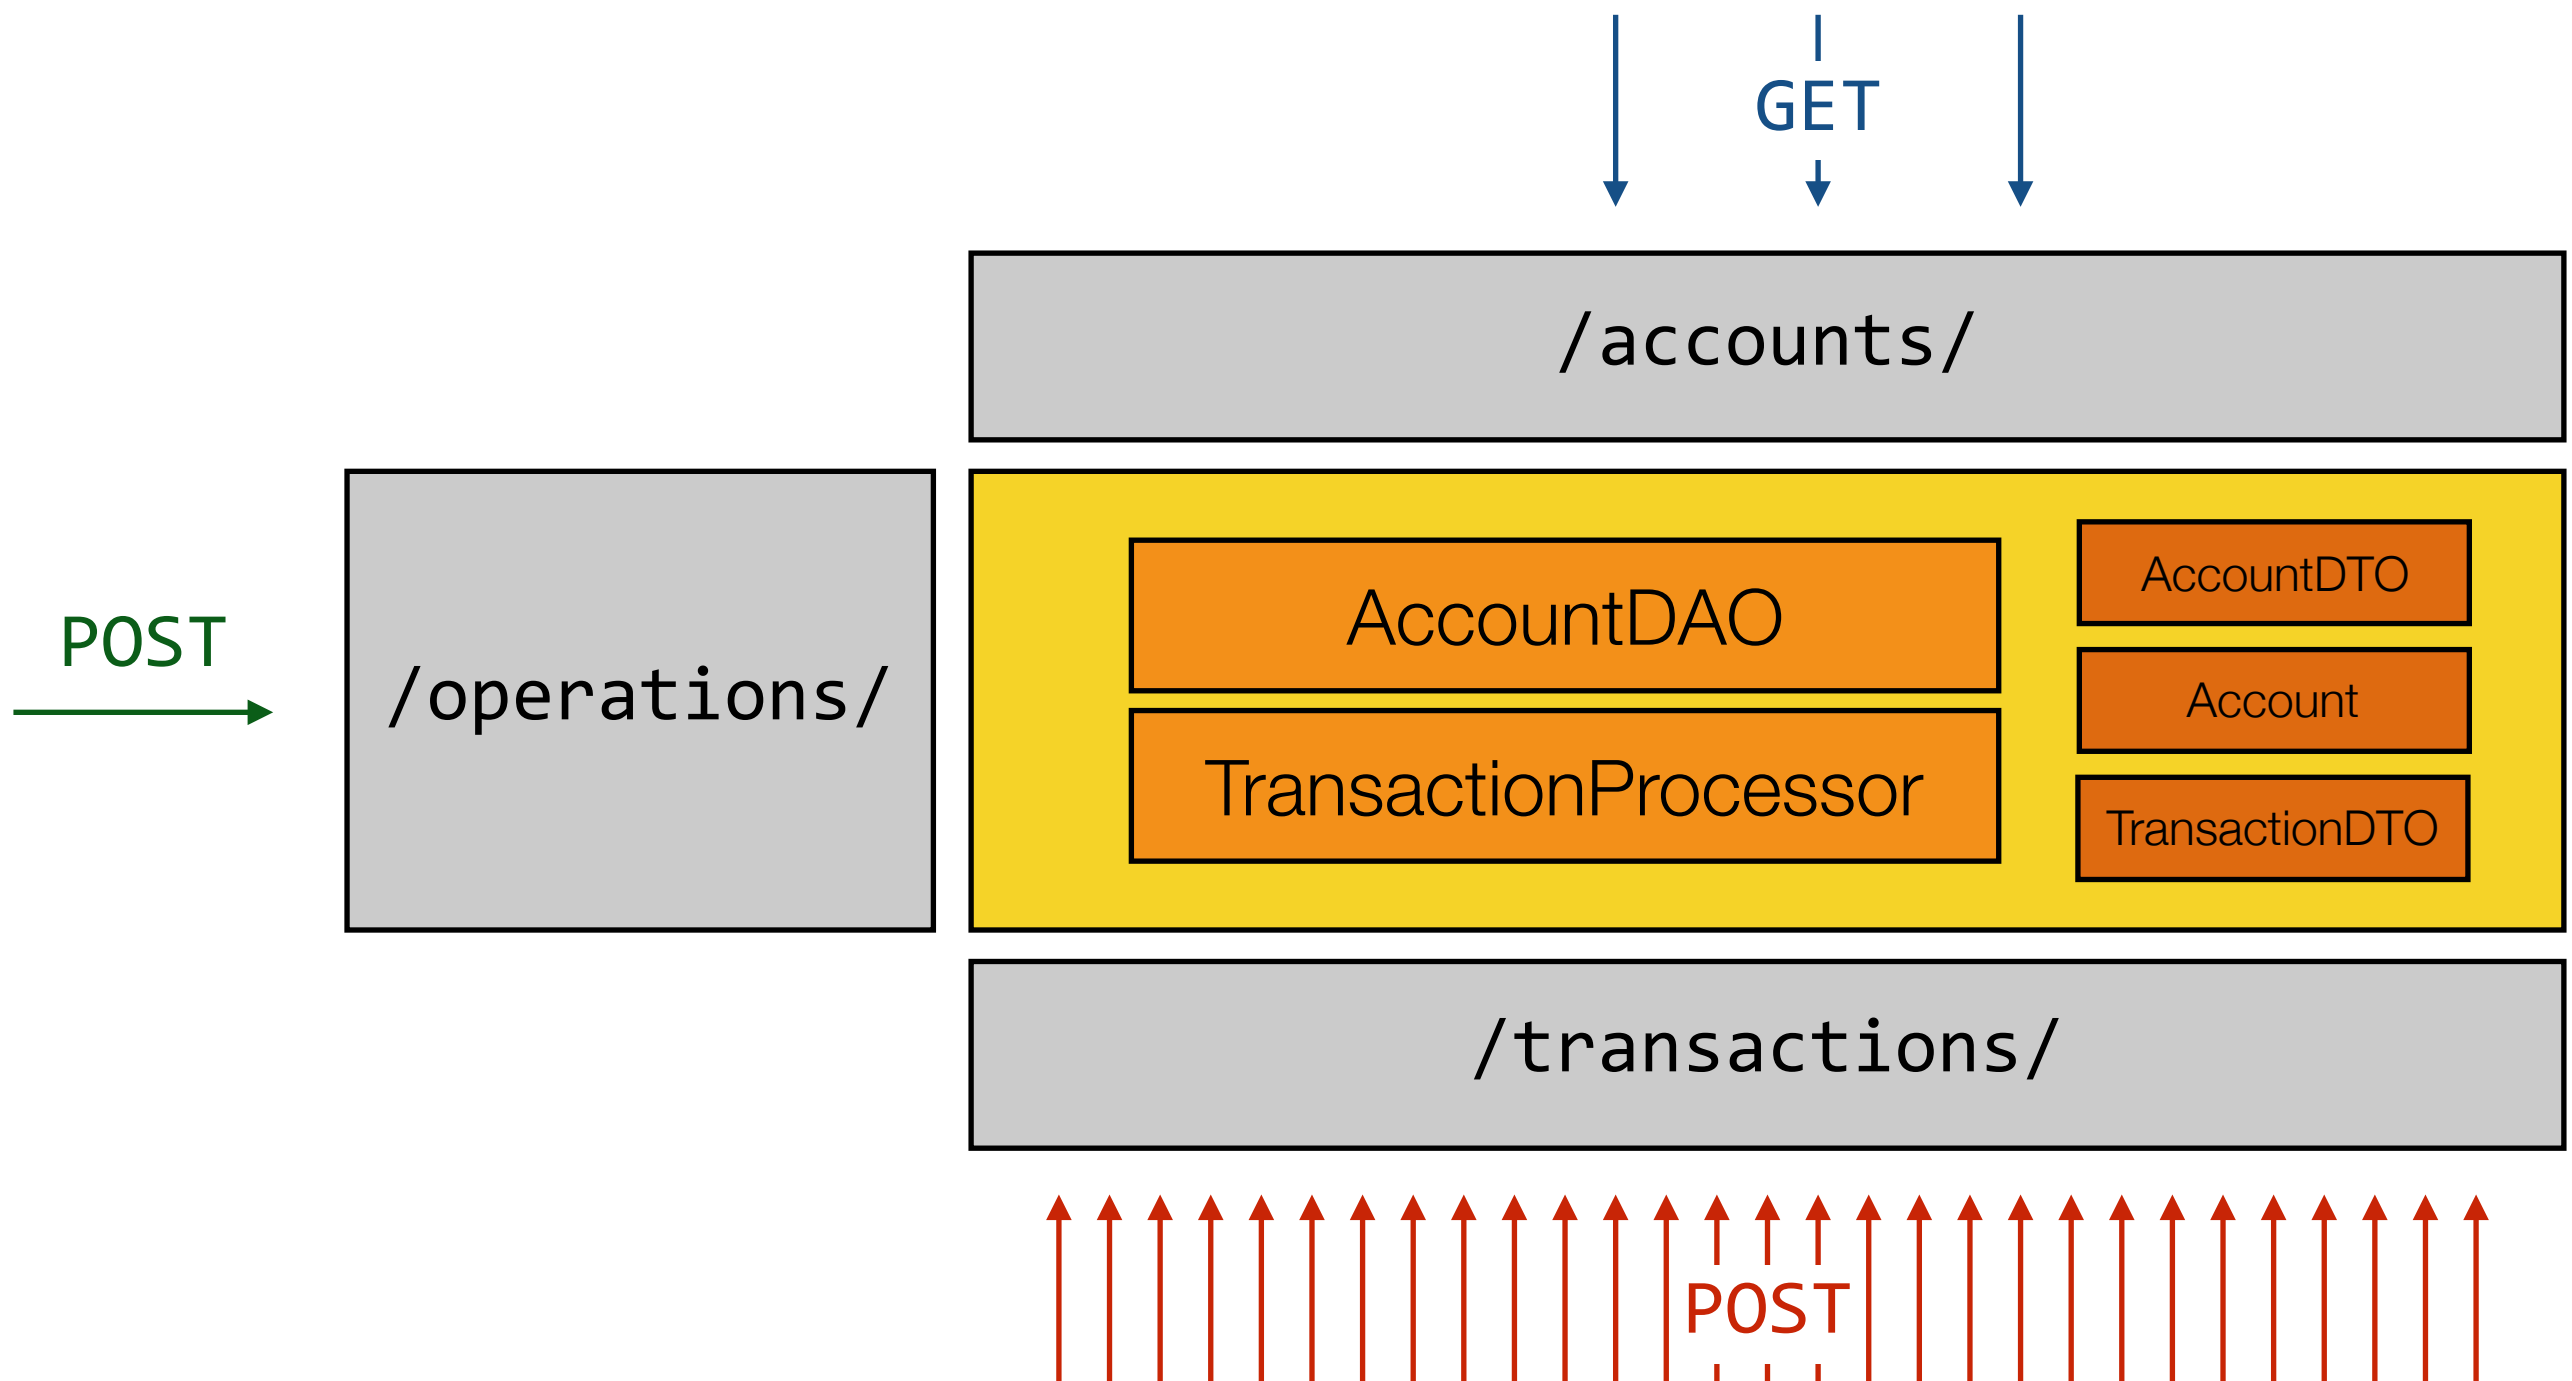
```
Account a = em.find(Account.class, id);  
em.lock(a, PESSIMISTIC_WRITE);
```

Be aware that we  
still have a risk of  
stale data here!

```
Account a = em.find(Account.class, id, PESSIMISTIC_WRITE);
```

# It works on my machine...

# System overview: REST APIs



# Concurrent creation & unique constraints



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
[www.heig-vd.ch](http://www.heig-vd.ch)

- In the system, we do not want to create accounts in advance.
- Instead, **we want to create them "on the fly"**: when we process a financial transaction, we check if the related account already exists:
  - If **no**, we create and initialize it.
  - If **yes**, we update it.
- **Let's try to implement this behavior!**

# Concurrent creation & unique constraints



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
www.heig-vd.ch

```
$ git clone git@github.com:SoftEng-HEIGVD/Teaching-HEIGVD-AMT-ConcurrentTransactions.git
```

```
$ git checkout step1-validating-on-the-fly-account-creation
```

Configure your **JDBC data source** (see `persistence.xml`): `jdbc/AMTDatabase`

After deploying the application, POST a number of transactions on `/api/transactions`. Then validate with GET `/accounts/`.

Check the Glassfish logs.

**Looks good!**

**Let's see how it runs...**



# Concurrent creation & unique constraints

The screenshot displays three concurrent REST client requests and their responses. The first two requests are for account 4242 with amounts 100 and 245. The third request is for account 1234 with amount 2048. The responses show the state of the accounts after each transaction, including the number of transactions and the holder's name.

**Request 1:** `POST http://localhost:8080/ConcurrentTransactionServer/api/transactions`  
Content-type: application/json  
Body: `{ "accountId": 4242, "amount": 100 }`

**Response 1:** `{ "id": 1234, "balance": 2048, "numberOfTransactions": 1, "holderName": "John Smith" }`

**Request 2:** `POST http://localhost:8080/ConcurrentTransactionServer/api/transactions`  
Content-type: application/json  
Body: `{ "accountId": 4242, "amount": 245 }`

**Response 2:** `{ "id": 4242, "balance": 245, "numberOfTransactions": 2, "holderName": "Frank Mueller" }`

**Request 3:** `POST http://localhost:8080/ConcurrentTransactionServer/api/transactions`  
Content-type: application/json  
Body: `{ "accountId": 1234, "amount": 2048 }`

**Response 3:** `{ "id": 1234, "balance": 2048, "numberOfTransactions": 1, "holderName": "John Smith" }`

Info: Received transaction for account: 4'242 100  
Info: \*\*\* Updating account: 4242 - 1  
Info: Received transaction for account: 4'242 245  
Info: \*\*\* Updating account: 4242 - 2  
Info: Received transaction for account: 1'234 2'048  
Info: \*\*\* Updating account: 1234 - 1

# Concurrent creation & unique constraints



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
[www.heig-vd.ch](http://www.heig-vd.ch)

Are we really **safe**?

```
$ git checkout step2-really-validating-on-the-fly-  
account-creation
```

**You now have 2 test projects:**

ConcurrentUpdateDemoClient (Java with JAX-RS client)

ConcurrentUpdateDemoClientNode (JavaScript)

**Let's see what happens “out-of-the-box”**

# The test clients

# General approach

---

- In order to **validate** that the server is working as expected, we need to:
  - **Generate some load** (simulate the activity of users, which translates into HTTP requests being sent to the API implementation).
  - Create and update a model, which **captures the expected state** of the domain model data at the end of the process (i.e. “if everything works well, this is how the ‘world’ should be).
  - **Keep track of errors reported by the server** (i.e. if the server tells us that it was unable to create a business object, we should not wrongly update the expect state). In other words, we have to make the difference between **known errors** and **silent bugs**.
  - At the end of the process, we have to **compare** the actual state on the server side with the model that we have built on the client side. If we see differences, then we have a problem (a silent bug).

# The Java test client

```
public class TestClient {  
    private void test() {  
        sendResetCommandToServer();  
  
        ExecutorService executor = Executors.newFixedThreadPool(numberOfConcurrentThreads);  
  
        final WebTarget target = client.target("http://localhost:8080/ConcurrentTransactionsServer/api").path("transactions");  
  
        for (int account = 1; account <= numberOfAccounts; account++) {  
            for (int transaction = 0; transaction < numberOfTransactionsPerAccount; transaction++) {  
                final int accountId = account;  
                Runnable task = new Runnable() {  
                    public void run() {  
                        TransactionDTO transaction = new TransactionDTO(accountId, 1);  
                        Response response = target.request().post(Entity.json(transaction));  
                        if (response.getStatus() < 200 || response.getStatus() >= 300) {  
                        } else {  
                            expectedState.logTransactionIntoAccount(transaction);  
                        }  
                    }  
                };  
                executor.execute(task);  
            }  
        }  
  
        try {  
            executor.shutdown();  
            executor.awaitTermination(1, TimeUnit.HOURS);  
            List<String> errors = validateExpectedAgainstActualState();  
        } catch (InterruptedException ex) {  
            Logger.getLogger(TestClient.class.getName()).log(Level.SEVERE, null, ex);  
        }  
    }  
}
```

**We reset the state on the server side (delete all accounts)**

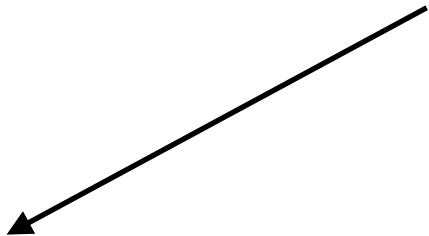
**In this version, the amount is 1 (but we could also generate a random value)**

**We only update the expected state (client side) if the server has told us that the transaction could be processed with 2xx HTTP status code (note that this works if the processing is synchronous)**

# The Java test client

In our application, the state is defined by a map of accounts. At the end of the process, we want to **check that the number of accounts** on the client and server sides are the same.

We also want to **check that the balance for every single account** is the same on the client and on the server side.



```
public class ExpectedState {

    private final Map<Long, AccountDTO> accounts = new HashMap<>();


    public synchronized void logTransactionIntoAccount(TransactionDTO transaction) {
        AccountDTO account = accounts.get(transaction.getAccountId());
        if (account == null) {
            account = new AccountDTO();
            account.setId(transaction.getAccountId());
            account.setNumberOfTransactions(0);
            account.setBalance(0);
            accounts.put(account.getId(), account);
        }
        account.setBalance(account.getBalance() + transaction.getAmount());
        account.setNumberOfTransactions(account.getNumberOfTransactions() + 1);
    }

    public Map<Long, AccountDTO> getAccounts() {
        return accounts;
    }

    ...
}
```

# The JavaScript test client


We use `async.js` to perform **3 main tasks** one after the other:  
**reset** the server state, **simulate** the users activity and **check** that  
the state on the client and server side is the same.



```
async.series([
  resetServerState,
  postTransactionRequestsInParallel,
  checkValues
], function(err, results) {
  console.log("\n\n=====");
  console.log("Summary");
  console.log("-----");
  //console.log(err);
  console.log(results);
});
```

# The JavaScript test client

Resetting the server state is easy, since we have a REST endpoint for that purpose.



```
function resetServerState( callback ) {

    console.log("\n\n=====");
    console.log("POSTing RESET command.");
    console.log("-----");

    client.post("http://localhost:8080/ConcurrentTransactionsServer/api/operations/
resetOperation", function(data, response) {
        console.log("RESET response status code: " + response.statusCode);
        callback(null, "The RESET operation has been processed (status code: " +
response.statusCode + ")");
    });

};
```



# The JavaScript test client


To simulate the activity of users, we use **async.js** once again. But this time, we execute multiple functions (to submit multiple requests) in **parallel**.

We only update the client-side state if the server has responded with a **2xx** status code.

```
function postTransactionRequestsInParallel(callback) {
  console.log("\n\n=====");
  console.log("POSTing transaction requests in parallel");
  console.log("-----");
  var numberOfUnsuccessfulResponses = 0;
  async.parallel(requests, function(err, results) {
    for (var i=0; i<results.length; i++) {
      if (results[i].response.statusCode < 200 || results[i].response.statusCode >= 300) {
        console.log("Result " + i + ": " + results[i].response.statusCode);
        numberOfUnsuccessfulResponses++;
      } else {
        logTransaction(processedStats, results[i].requestData.data);
      }
    }
    callback(null, results.length + " transaction POSTs have been sent. " +
      numberOfUnsuccessfulResponses + " have failed.");
  });
}
```

# The JavaScript test client

When we use `async.js` to execute functions in parallel, we must provide an array of functions. We prepare it in advance.



```
var requests = [];  
for (var account=1; account<=numberOfAccounts; account++) {  
    for (var transaction=0; transaction<numberOfTransactionsPerAccount; transaction++) {  
        requests.push(  
            getTransactionPOSTRequestFunction(account)  
        );  
    }  
};
```

```
function getTransactionPOSTRequestFunction(accountId) {  
    return function(callback) {  
        var requestData = { ... }  
    };  
    requestData.data.amount = Math.floor((Math.random() * 200) - 50);  
    logTransaction(submittedStats, requestData.data);  
    client.post("http://localhost:8080/.../api/transactions", requestData, function(data, response) {  
        var error = null;  
        var result = { requestData: requestData, data: data, response: response };  
        callback(error, result);  
    });  
}
```

We also **keep track** of all the transactions that we have submitted (some may be rejected by the server). This is an additional feature compared to the Java version.

# Running the tests...

# ConcurrentUpdateDemoClientNode

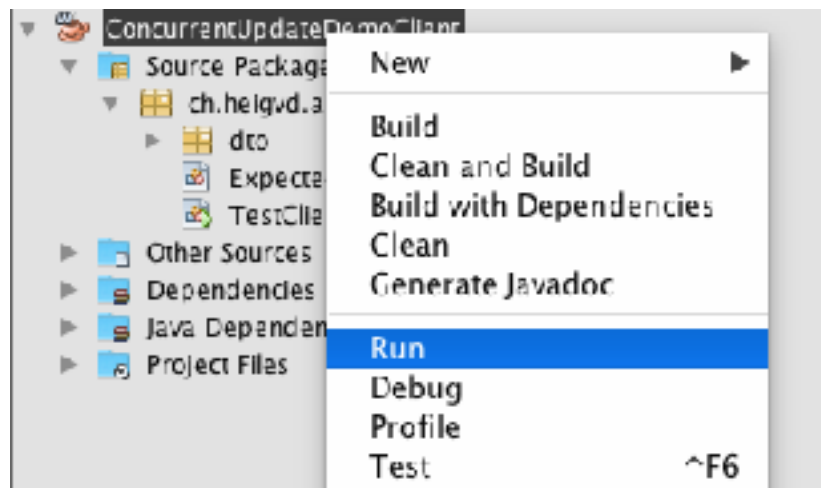
```
$ npm install
$ node client.js
```

```
=====
Comparing client-side and server-side stats
-----
Number of accounts on the client side: 10
Number of accounts on the server side: 10
```

```
=====
Summary
-----
[ 'The RESET operation has been processed (status code: 204)',
  '200 transaction POSTs have been sent. 0 have failed.',
  'The client side and server side values have been compared. Number of corrupted accounts: 0' ]
```

```
Info: Received transaction for account: 1 33
Info: *** Updating account: 1 - 1
Info: Received transaction for account: 1 74
Info: *** Updating account: 1 - 2
Info: Received transaction for account: 1 85
Info: *** Updating account: 1 - 3
Info: Received transaction for account: 1 1
Info: *** Updating account: 1 - 4
Info: Received transaction for account: 1 118
Info: *** Updating account: 1 - 5
Info: Received transaction for account: 1 -11
Info: *** Updating account: 1 - 6
Info: Received transaction for account: 1 61
Info: *** Updating account: 1 - 7
Info: Received transaction for account: 1 -3
Info: *** Updating account: 1 - 8
Info: Received transaction for account: 1 126
Info: *** Updating account: 1 - 9
Info: Received transaction for account: 1 -28
Info: *** Updating account: 1 - 10
```

# ConcurrentUpdateDemoClient



```
10:50:54 INFO Expected vs actual number of transactions for account 18: 20/20
10:50:54 INFO Expected vs actual balance for account 18: 20/20
10:50:54 INFO Expected vs actual number of transactions for account 19: 20/20
10:50:54 INFO Expected vs actual balance for account 19: 20/20
10:50:54 INFO Expected vs actual number of transactions for account 20: 20/20
10:50:54 INFO Expected vs actual balance for account 20: 20/20
10:50:54 INFO Errors: []
10:50:54 INFO Done.
```

```
Info: *** Updating account: 20 - 1
Info: Received transaction for account: 20 1
Info: *** Updating account: 20 - 2
Info: Received transaction for account: 20 1
Info: *** Updating account: 20 - 3
Info: Received transaction for account: 20 1
Info: *** Updating account: 20 - 4
Info: Received transaction for account: 20 1
Info: *** Updating account: 20 - 5
Info: Received transaction for account: 20 1
Info: *** Updating account: 20 - 6
Info: Received transaction for account: 20 1
Info: *** Updating account: 20 - 7
```

# Concurrent creation & unique constraints



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
[www.heig-vd.ch](http://www.heig-vd.ch)

Still looks good... Are we really **safe**?

**Change the experiment parameters, so that we have  
concurrent requests!**

There are parameters in the Java and the JavaScript test client

# ConcurrentUpdateDemoClientNode

```
$ node client.js
```

```
Result 162: 500  
Result 181: 500
```

Some POST requests fail (the client is aware of a problem)

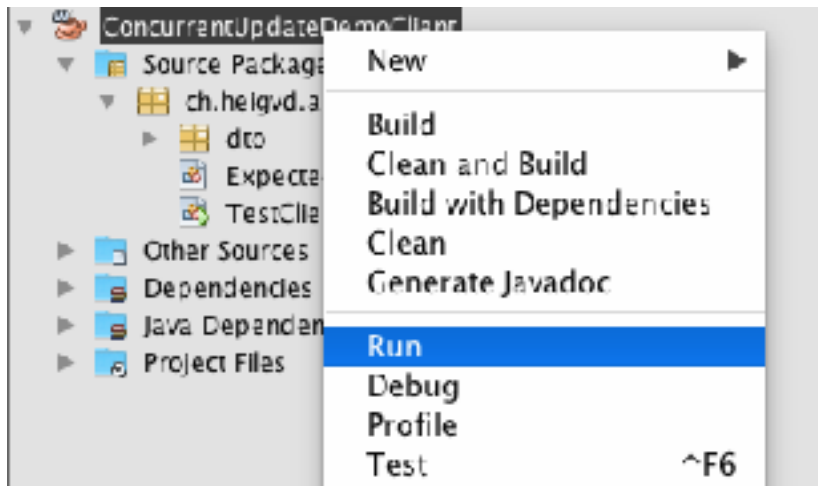
```
=====
Comparing client-side and server-side stats
-----
```

```
Number of accounts on the client side: 10
Number of accounts on the server side: 10
Account 1 --> Server/Client balance: 276/908 X
Account 2 --> Server/Client balance: 573/1161 X
Account 3 --> Server/Client balance: 478/1007 X
Account 4 --> Server/Client balance: 280/532 X
Account 5 --> Server/Client balance: 612/923 X
Account 6 --> Server/Client balance: 192/942 X
Account 7 --> Server/Client balance: 342/722 X
Account 8 --> Server/Client balance: 555/800 X
Account 9 --> Server/Client balance: 354/1107 X
Account 10 --> Server/Client balance: 407/1264 X
```

Worse: money has  
vanished without anyone  
being aware of it!

```
Caused by: javax.persistence.PersistenceException: Exception [EclipseLink-4002] (Eclipse Persistence Services - 2.5.2.v20140319-9ad6abd):
org.eclipse.persistence.exceptions.DatabaseException
Internal Exception: com.mysql.jdbc.exceptions.jdbc4.MySQLIntegrityConstraintViolationException: Duplicate entry '3' for key 'PRIMARY'
Error Code: 1062
Call: INSERT INTO ACCOUNT (ID, BALANCE, HOLDERNAME, NUMBEROFTRANSACTIONS) VALUES (?, ?, ?, ?)
bind => [4 parameters bound]
Query: InsertObjectQuery(ch.heigvd.amt.demo.model.Account@7512b0e3)
at org.eclipse.persistence.internal.jpa.EntityManagerImpl.flush(EntityManagerImpl.java:868)
at com.sun.enterprise.container.common.impl.EntityManagerWrapper.flush(EntityManagerWrapper.java:437)
at ch.heigvd.amt.demo.services.dao.AccountDAO.create(AccountDAO.java:26)
```

# ConcurrentUpdateDemoClient



```
tasks have been submitted to the executor and will be processed by 5 concurrent threads.  
server was not able to process the transaction: 500 Internal Server Error  
server was not able to process the transaction: 500 Internal Server Error  
server was not able to process the transaction: 500 Internal Server Error  
server was not able to process the transaction: 500 Internal Server Error  
server was not able to process the transaction: 500 Internal Server Error  
server was not able to process the transaction: 500 Internal Server Error
```

```
11:04:37 INFO Expected vs actual number of transactions for account 15: 20/5  
11:04:37 INFO Expected vs actual balance for account 15: 20/5  
11:04:37 INFO Expected vs actual number of transactions for account 16: 20/5  
11:04:37 INFO Expected vs actual balance for account 16: 20/5  
11:04:37 INFO Expected vs actual number of transactions for account 17: 20/5  
11:04:37 INFO Expected vs actual balance for account 17: 20/5  
11:04:37 INFO Expected vs actual number of transactions for account 18: 20/5  
11:04:37 INFO Expected vs actual balance for account 18: 20/5  
11:04:37 INFO Expected vs actual number of transactions for account 19: 20/5  
11:04:37 INFO Expected vs actual balance for account 19: 20/5  
11:04:37 INFO Expected vs actual number of transactions for account 20: 20/6  
11:04:37 INFO Expected vs actual balance for account 20: 20/6  
11:04:37 INFO Errors: [The number of transactions for account 1 is not the one expected: 7 vs 17, The balance  
for account 1 is not the one expected: 7.0 vs 17.0
```

```
Caused by: javax.persistence.PersistenceException: Exception [EclipseLink-4002] (Eclipse Persistence Services - 2.5.2.v20140319-9ad6abd):  
org.eclipse.persistence.exceptions.DatabaseException  
Internal Exception: com.mysql.jdbc.exceptions.jdbc4.MySQLIntegrityConstraintViolationException: Duplicate entry '1' for key 'PRIMARY'  
Error Code: 1062  
Call: INSERT INTO ACCOUNT (ID, BALANCE, HOLDERNAME, NUMBEROFTRANSACTIONS) VALUES (?, ?, ?, ?)  
bind => [4 parameters bound]  
Query: InsertObjectQuery(ch.heigvd.amt.demo.model.Account@178ebe08)  
at org.eclipse.persistence.internal.jpa.EntityManagerImpl.flush(EntityManagerImpl.java:868)  
at com.sun.enterprise.container.common.impl.EntityManagerWrapper.flush(EntityManagerWrapper.java:437)  
at ch.heigvd.amt.demo.services.dao.AccountDAO.create(AccountDAO.java:26)
```



# The account creation issue

# What is the **account creation** problem?



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
[www.heig-vd.ch](http://www.heig-vd.ch)

```
@Override
public void createAccountIfNotExists(long id) {
    Account account = accountDAO.findById(id);
    if (account == null) {
        account = new Account();
        account.setId(id);
        account.setBalance(0);
        account.setNumberOfTransactions(0);
        account.setHolderName(generateRandomHolderName());
        accountDAO.create(account);
    }
}
```

## Thread T1 on EJB 1

```
Account account = accountDAO.findById(id);
if (account == null) {
    account = new Account();
    account.setId(id);
    account.setBalance(0);
    account.setNumberOfTransactions(0);
    account.setHolderName(generateRandomHolderName());
```

```
    accountDAO.create(account);
}
```



## Thread T2 on EJB2

```
Account account = accountDAO.findById(id);
if (account == null) {
    account = new Account();
    account.setId(id);
    account.setBalance(0);
    account.setNumberOfTransactions(0);
    account.setHolderName(generateRandomHolderName());
    accountDAO.create(account);
}
```

# Fixing the problem: approach 1 (new tx)

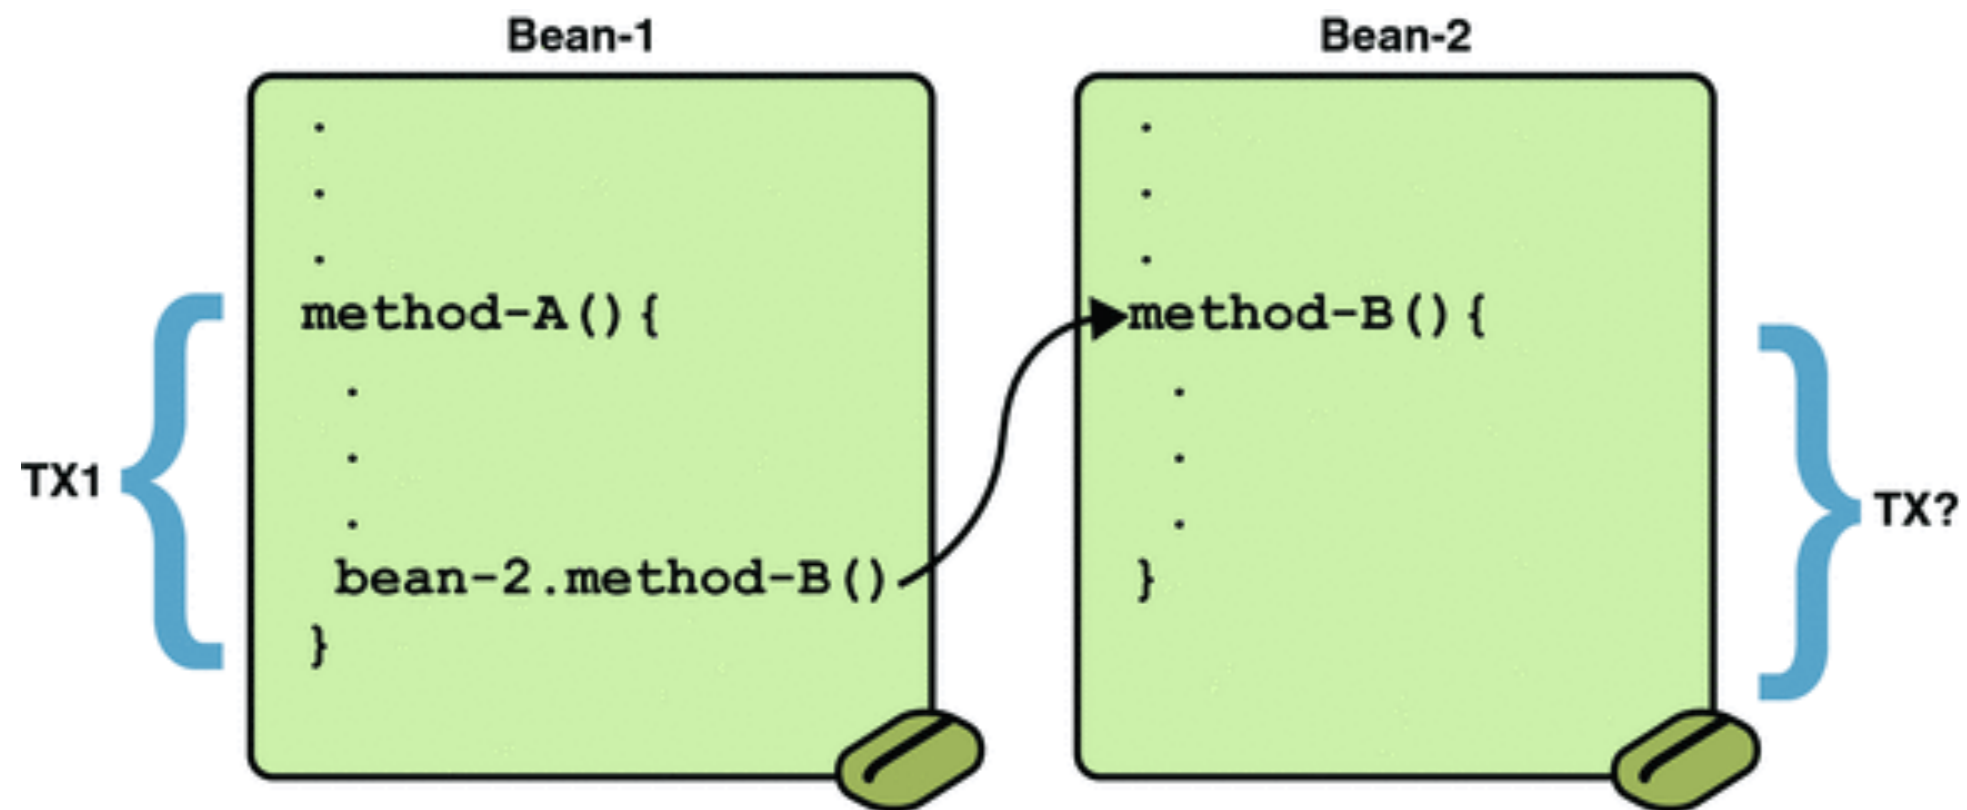


HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
[www.heig-vd.ch](http://www.heig-vd.ch)

- In a previous lecture, we have seen that it is possible:
  - to divide one “use case” into multiple sub-transactions
  - to decide whether all sub-transactions should be rolled back or only some of them in the case of errors
- We have seen that there is a special annotation (`@TransactionAttribute`) for specifying the behavior (by default, the container rolls back everything).

```
$ git checkout step3-fix-account-creation-with-try-catch
```

# Transaction Scope



<http://java.sun.com/javase/5/docs/tutorial/doc/bncij.html>


# Transaction Scope

```
@TransactionalAttribute(NOT_SUPPORTED)
@Stateless
public class TransactionBean implements
Transaction {
...
@TransactionalAttribute(REQUIRES_NEW)
public void firstMethod() {...}

@TransactionalAttribute(REQUIRED)
public void secondMethod() {...}

public void thirdMethod() {...}

public void fourthMethod() {...}
}
```



Transaction Attribute	Client's Transaction	Business Method's Transaction
Required	None	T2
	T1	T1
RequiresNew	None	T2
	T1	T2
Mandatory	None	error
	T1	T1
NotSupported	None	None
	T1	None
Supports	None	None
	T1	T1
Never	None	None
	T1	Error

# Fixing the problem: approach 1 (new tx)

If we want to capture failed transactions, we need to go via the container

```
@Stateless
public class TransactionProcessor implements TransactionProcessorLocal {

    private static final Logger LOG = Logger.getLogger(TransactionProcessor.class.getName());

    @EJB
    AccountDAOLocal accountDAO;

    @EJB
    TransactionProcessorLocal selfViaContainer;

    @Override
    public void processTransaction(TransactionDTO transaction) {
        try {
            selfViaContainer.createAccountIfNotExists(transaction.getAccountId());
        } catch (Exception e) {
            LOG.info("*** Maybe a DUPLICATE KEY that would not be a real problem..." + e.getMessage());
        }
        ...
    }

    @TransactionAttribute(TransactionAttributeType.REQUIRES_NEW)
    public void createAccountIfNotExists(long id) {
        Account account = accountDAO.findById(id);
        if (account == null) {
            account = new Account();
            account.setId(id);
            account.setBalance(0);
            account.setNumberOfTransactions(0);
            account.setHolderName(generateRandomHolderName());
            accountDAO.create(account);
        }
    }
}
```

If an exception occurs in this block, we don't want to rollback everything!

# ConcurrentUpdateDemoClientNode

```
$ node client.js
```

```
=====
Summary
-----
[ 'The RESET operation has been processed (status code: 204)',
  '200 transaction POSTs have been sent. 0 have failed.',
  'The client side and server side values have been compared. Number of corrupted accounts: 10' ]
```

We have resolved one issue: the client does not receive any error when the first two financial transactions for one account are sent simultaneously.

However, we still have a problem with data corruption (unrelated to account creation).

We also have ugly stack traces in our logs

```
Caused by: javax.persistence.PersistenceException: Exception [EclipseLink-4002] (Eclipse Persistence Services - 2.5.2.v20140319-9ad6abd):
org.eclipse.persistence.exceptions.DatabaseException
Internal Exception: com.mysql.jdbc.exceptions.jdbc4.MySQLIntegrityConstraintViolationException: Duplicate entry '3' for key 'PRIMARY'
Error Code: 1062
Call: INSERT INTO ACCOUNT (ID, BALANCE, HOLDERNAME, NUMBEROFTRANSACTIONS) VALUES (?, ?, ?, ?)
      bind => [4 parameters bound]
Query: InsertObjectQuery(ch.heigvd.amt.demo.model.Account@7512b0e3)
      at org.eclipse.persistence.internal.jpa.EntityManagerImpl.flush(EntityManagerImpl.java:868)
      at com.sun.enterprise.container.common.impl.EntityManagerWrapper.flush(EntityManagerWrapper.java:437)
      at ch.heigvd.amt.demo.services.dao.AccountDAO.create(AccountDAO.java:26)
```

# Fixing the problem: approach 2 (upsert)



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
[www.heig-vd.ch](http://www.heig-vd.ch)

- Many databases support a special type of operation, often called an “upsert”
- With this operation, you can specify that when you can update a record if it already exists in the database, or create it if does not exist yet.
- MySQL supports this feature with the **INSERT ... ON DUPLICATE KEY UPDATE syntax**

```
$ git checkout step4-fix-account-creation-with-upsert
```



# Fixing the problem: approach 2 (upsert)



HAUTE ÉCOLE  
D'INGÉNIERIE ET DE GESTION  
DU CANTON DE VAUD  
www.heig-vd.ch

```
@Entity
@NamedQueries({
    @NamedQuery(name="Account.findAll", query="SELECT a FROM Account a"),
    @NamedQuery(name="Account.deleteAll", query="DELETE FROM Account")
})
@NamedNativeQuery(name = "Account.upsert", query = "INSERT INTO Account (ID, HOLDERNAME, BALANCE, NUMBEROFTRANSACTIONS) VALUES
(?1, ?2, ?3, ?4) ON DUPLICATE KEY UPDATE BALANCE=BALANCE+?4, NUMBEROFTRANSACTIONS=NUMBEROFTRANSACTIONS+0")
public class Account { ... }
```

This is a **proprietary** feature  
provided by MySQL

```
@Stateless
public class TransactionProcessor implements TransactionProcessorLocal {

    private static final Logger LOG = Logger.getLogger(TransactionProcessor.class.getName());

    @EJB
    AccountDAOLocal accountDAO;

    public void createAccountIfNotExists(long id) {
        @Override
        public void createAccountIfNotExists(long id) {
            Query query = em.createNamedQuery("Account.upsert");
            query.setParameter(1, id);
            query.setParameter(2, generateRandomHolderName());
            query.setParameter(3, 0);
            query.setParameter(4, 0);
            long result = query.executeUpdate();
        }
    }
}
```

# ConcurrentUpdateDemoClientNode

```
$ node client.js
```

```
=====
Comparing client-side and server-side stats
-----
Number of accounts on the client side: 20
Number of accounts on the server side: 20

=====
Summary
-----
[ 'The RESET operation has been processed (status code: 204)',
  '800 transaction POSTs have been sent. 0 have failed.',
  'The client side and server side values have been compared. Number of corrupted accounts: 0' ]
```

We don't have any 500 response sent to the client (no problem with duplicate accounts)

As an additional benefit, we don't have any data corruption! That is because the special MySQL requests locks the row in the database.

We have also got rid of the exceptions!

```
Info: Received transaction for account: 20 104
Info: *** Updating account: 20 - 22
Info: Received transaction for account: 20 21
Info: *** Updating account: 20 - 23
Info: Received transaction for account: 20 46
Info: *** Updating account: 20 - 24
Info: *** Updating account: 20 - 25
Info: Received transaction for account: 20 143
Info: *** Updating account: 20 - 26
Info: Received transaction for account: 20 74
Info: *** Updating account: 20 - 27
Info: Received transaction for account: 20 12
```

# The data corruption issue

# What is the **data corruption** problem?

```
public void processTransaction(TransactionDTO transaction) {  
    try {  
        selfViaContainer.createAccountIfNotExists(transaction.getAccountId());  
    } catch (Exception e) {  
        LOG.info("*** Maybe a DUPLICATE KEY that would not be a real problem..." + e.getMessage());  
    }  
  
    Account account = accountDAO.findById(transaction.getAccountId());  
    double bal = account.getBalance();  
    bal = bal + transaction.getAmount();  
    account.setBalance(bal);  
    account.setNumberOfTransactions(account.getNumberOfTransactions() + 1);  
}
```

Thread T1 on EJB 1

Thread T2 on EJB2

```
Account account = accountDAO.findById(transaction.getAccountId());  
double bal = account.getBalance();
```



```
bal = bal + transaction.getAmount();  
account.setBalance(bal);  
account.setNumberOfTransactions(account.getNumberOfTransactions() + 1);
```

```
Account account = accountDAO.findById(transaction.getAccountId());  
double bal = account.getBalance();  
bal = bal + transaction.getAmount();  
account.setBalance(bal);  
account.setNumberOfTransactions(account.getNumberOfTransactions() + 1);
```

# Optimistic vs Pessimistic Locking

- To fix this issue, we have the choice between a pessimistic and an optimistic locking strategy:
  - If we believe that there is a high probability to have a conflict (we are pessimistic), then we should lock the record before modifying it (the other transaction will have to wait that we release it).
  - If we believe that there is a little probability to have a conflict, then we can look at the version number of the record when we read it, check that it is still the same and increment it when we update the record.
  - If someone has modified the record in the meantime, then the version number will have changed and we will be aware of the issue.
- JPA provides support for both pessimistic and optimistic locking strategies.
- Pessimistic Locking has a performance cost (and may introduce deadlocks). Optimistic locking may require some extra work (dealing with exceptions).

# Pessimistic locking solution

```
$ git checkout step5-fix-account-creation-with-try-  
catch-pessimistic-lock
```

```
@Stateless  
public class AccountDAO implements AccountDAOLocal {  
  
    ...  
    @Override  
    public Account findByIdForUpdate(long id) {  
        return em.find(Account.class, id, LockModeType.PESSIMISTIC_WRITE);  
    }  
    ...  
}
```

```
@Stateless  
public class TransactionProcessor implements TransactionProcessorLocal {  
    ...  
    @Override  
    public void processTransaction(TransactionDTO transaction) {  
        ...  
        Account account = accountDAO.findByIdForUpdate(transaction.getAccountId());  
        ...  
    }  
}
```

# Optimistic locking solution

---

```
$ git checkout step6-fix-account-creation-with-try-  
catch-optimistic-lock
```

```
@Entity  
public class Account {  
  
    @Id  
    private long id;  
  
    @Version  
    private long version;  
}
```